

Homework 12 Mini Report

The time it took the CuPy parallelized matrix population and operations to complete was 4.401585 seconds. For a similar operation, Numpy took 58.65332 seconds to finish. Using CuPy for a 30000^2 matrix resulted in an over 13x speedup as compared to Numpy. This is likely because of the ability for CuPy to rapidly populate the K matrix and f vector with it comparatively vast number of workers. Additionally, populating the matrix and vector was a menial task, not requiring much power. There may also have been parallelism-based time saves with CuPy's `linalg.solve()` function that Numpy may not have been able to exploit due to a more limited number of workers.